

# Operating Systems

Steven Hand

*Michaelmas Term 2010*

12 lectures for CST IA

# Course Aims

---

- This course aims to:
  - explain the structure and functions of an operating system,
  - illustrate key operating system aspects by concrete example, and
  - prepare you for future courses. . .
- At the end of the course you should be able to:
  - compare and contrast CPU scheduling algorithms
  - explain the following: process, address space, file.
  - distinguish paged and segmented virtual memory.
  - discuss the relative merits of Unix and NT. . .

# Course Outline

---

- Introduction to Operating Systems.
- Processes & Scheduling.
- Memory Management.
- I/O & Device Management.
- Protection.
- Filing Systems.
- Case Study: Unix.
- Case Study: Windows NT.

## Recommended Reading

---

- *Concurrent Systems or Operating Systems*  
Bacon J [ and Harris T ], Addison Wesley 1997 [2003]
- *Operating Systems Concepts (5th Ed.)*  
Silberschatz A, Peterson J and Galvin P, Addison Wesley 1998.
- *The Design and Implementation of the 4.3BSD UNIX Operating System*  
Leffler S J, Addison Wesley 1989
- *Inside Windows 2000 (3rd Ed) or Windows Internals (4th Ed)*  
Solomon D and Russinovich M, Microsoft Press 2000 [2005]


# What is an Operating System?

---

- A program which controls the execution of all other programs (applications).
- Acts as an intermediary between the user(s) and the computer.
- Objectives:
  - convenience,
  - efficiency,
  - extensibility.
- Similar to a government. . .

# An Abstract View

---



- The Operating System (OS):
  - controls all execution.
  - multiplexes resources between applications.
  - abstracts away from complexity.
- Typically also have some *libraries* and some *tools* provided with OS.
- Are these part of the OS? Is IE a tool?
  - no-one can agree. . .
- For us, the OS  $\approx$  the *kernel*.

# In The Beginning. . .

---

- 1949: First stored-program machine (EDSAC)
- to ~ 1955: “Open Shop”.
  - large machines with vacuum tubes.
  - I/O by paper tape / punch cards.
  - user = programmer = operator.
- To reduce cost, hire an *operator*:
  - programmers write programs and submit tape/cards to operator.
  - operator feeds cards, collects output from printer.
- Management like it.
- Programmers hate it.
- Operators hate it.

⇒ need something better.

# Batch Systems


---

- Introduction of tape drives allow *batching* of jobs:
  - programmers put jobs on cards as before.
  - all cards read onto a tape.
  - operator carries input tape to computer.
  - results written to output tape.
  - output tape taken to printer.
- Computer now has a *resident monitor*:
  - initially control is in monitor.
  - monitor reads job and transfer control.
  - at end of job, control transfers back to monitor.
- Even better: *spooling systems*.
  - use interrupt driven I/O.
  - use magnetic disk to cache input tape.
  - fire operator.
- Monitor now *schedules* jobs. . .



# Multi-Programming

---



- Use memory to cache jobs from disk  $\Rightarrow$  more than one job active simultaneously.
- Two stage scheduling:
  1. select jobs to load: *job scheduling*.
  2. select resident job to run: *CPU scheduling*.
- Users want more interaction  $\Rightarrow$  *time-sharing*:
- e.g. CTSS, TSO, Unix, VMS, Windows NT. . .


# Today and Tomorrow

---

- Single user systems: cheap and cheerful.
  - personal computers.
  - no other users  $\Rightarrow$  ignore protection.
  - e.g. DOS, Windows, Win 95/98, . . .
- RT Systems: power is nothing without control.
  - hard-real time: nuclear reactor safety monitor.
  - soft-real time: mp3 player.
- Parallel Processing: the need for speed.
  - SMP: 2–8 processors in a box.
  - MIMD: super-computing.
- Distributed computing: global processing?
  - Java: the network is the computer.
  - Clustering: the network is the bus.
  - CORBA: the computer is the network.
  - .NET: the network is an enabling framework. . .

# Monolithic Operating Systems

---



- Oldest kind of OS structure (“modern” examples are DOS, original MacOS)
- Problem: applications can e.g.
  - trash OS software.
  - trash another application.
  - hoard CPU time.
  - abuse I/O devices.
  - etc. . .
- No good for fault containment (or multi-user).
- Need a better solution. . .

# Dual-Mode Operation


---

- Want to stop buggy (or malicious) program from doing bad things.

⇒ provide *hardware* support to distinguish between (at least) two different modes of operation:

1. *User Mode* : when executing on behalf of a user (i.e. application programs).
2. *Kernel Mode* : when executing on behalf of the operating system.

- Hardware contains a mode-bit, e.g. 0 means kernel, 1 means user.




- Make certain machine instructions only possible in kernel mode. . .

# Protecting I/O & Memory

---

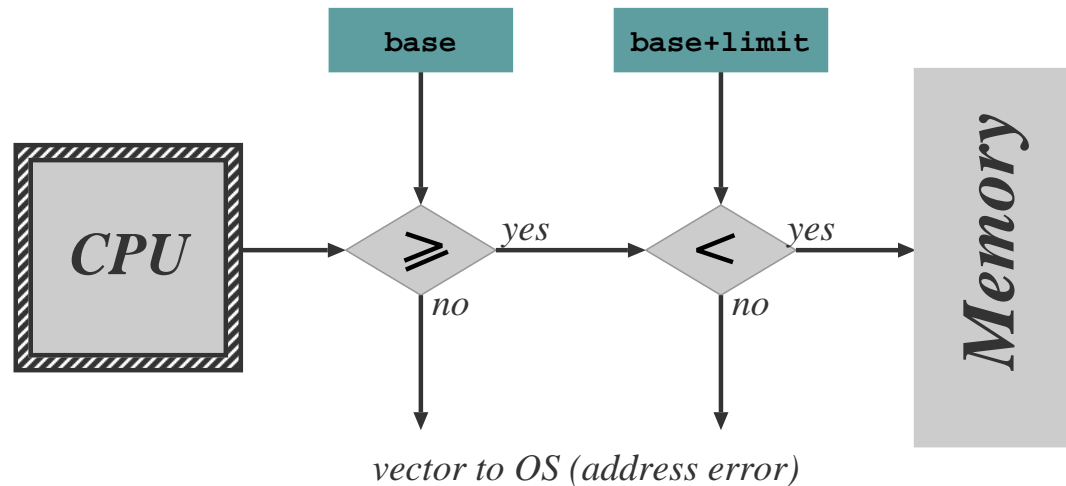
- First try: make I/O instructions privileged.
  - applications can't mask interrupts.
  - applications can't control I/O devices.
- But:
  1. Application can rewrite interrupt vectors.
  2. Some devices accessed via *memory*
- Hence need to protect memory also, e.g. define *base* and *limit* for each program:



- Accesses outside allowed range are protected.

# Memory Protection Hardware

---



- Hardware checks every memory reference.
- Access out of range  $\Rightarrow$  vector into operating system (just as for an interrupt).
- Only allow *update* of base and limit registers in kernel mode.
- Typically disable memory protection in kernel mode (although a bad idea).
- In reality, more complex protection h/w used:
  - main schemes are *segmentation* and *paging*
  - (covered later on in course)

# Protecting the CPU

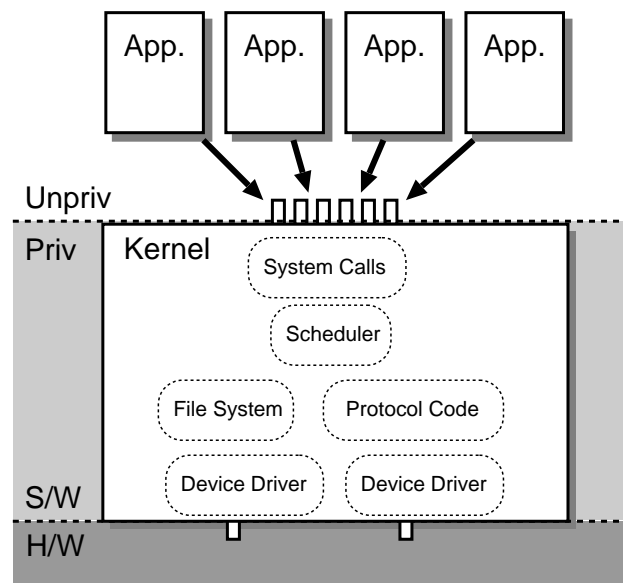
---

- Need to ensure that the OS stays in control.
  - i.e. need to prevent any a malicious or badly-written application from ‘hogging’ the CPU the whole time.

⇒ use a *timer* device.
- Usually use a *countdown* timer, e.g.
  1. set timer to initial value (e.g. 0xFFFF).
  2. every *tick* (e.g.  $1\mu s$ ), timer decrements value.
  3. when value hits zero, interrupt.
- (Modern timers have programmable tick rate.)
- Hence OS gets to run periodically and do its stuff.
- Need to ensure only OS can load timer, and that interrupt cannot be masked.
  - use same scheme as for other devices.
  - (viz. privileged instructions, memory protection)
- Same scheme can be used to implement time-sharing (more on this later).

# Kernel-Based Operating Systems

---




- Applications can't do I/O due to protection  
⇒ operating system does it on their behalf.
- Need secure way for application to invoke operating system:  
⇒ require a special (unprivileged) instruction to allow transition from user to kernel mode.
- Generally called a *software interrupt* since operates similarly to a real (hardware) interrupt. . .
- Set of OS services accessible via software interrupt mechanism called *system calls*.



# Microkernel Operating Systems

---



- Alternative structure:
  - push some OS services into *servers*.
  - servers may be privileged (i.e. operate in kernel mode).
- Increases both *modularity* and *extensibility*.
- Still access kernel via system calls, but need new way to access servers:
  - ⇒ interprocess communication (IPC) schemes.

# Kernels versus Microkernels

---

So why isn't everything a microkernel?

- Lots of IPC adds overhead  
⇒ microkernels usually perform less well.
- Microkernel implementation sometimes tricky: need to worry about concurrency and synchronisation.
- Microkernels often end up with redundant copies of OS data structures.

Hence today most common operating systems blur the distinction between kernel and microkernel.

- e.g. linux is a “kernel”, but has kernel modules and certain servers.
- e.g. Windows NT was originally microkernel (3.5), but now (4.0 onwards) pushed lots back into kernel for performance.
- Still not clear what the best OS structure is, or how much it really matters. . .

# Operating System Functions

---

- Regardless of structure, OS needs to *securely multiplex resources*:
  1. protect applications from each other, yet
  2. share physical resources between them.
- Also usually want to *abstract* away from grungy hardware, i.e. OS provides a *virtual machine*:
  - share CPU (in time) and provide each app with a **virtual processor**,
  - allocate and protect memory, and provide applications with their own **virtual address space**,
  - present a set of (relatively) hardware independent **virtual devices**,
  - divide up storage space by using **filing systems**, and
  - do all this within the context of a **security framework**.
- Remainder of this part of the course will look at each of the above areas in turn. . .

# Process Concept

---

- From a user's point of view, the operating system is there to execute programs:
  - on batch system, refer to *jobs*
  - on interactive system, refer to *processes*
  - (we'll use both terms fairly interchangeably)
- Process  $\neq$  Program:
  - a program is *static*, while a process is *dynamic*
  - in fact, a process  $\triangleq$  “a program in execution”
- (Note: “program” here is pretty low level, i.e. native machine code or *executable*)
- Process includes:
  1. program counter
  2. stack
  3. data section
- Processes execute on *virtual processors*